RANKED, NOT CITED

Your page ranks on Google. Why doesn’t AI cite it?

150

characters: the most of a cited page that Claude’s web search returns as cited text

One documented reason: AI reads more than it quotes.

Claude Platform Docs, web search tool: “cited_text: Up to 150 characters of the cited content.” Page has no date; checked 4 Nov 2026.

Every engine sets its own limits

CLAUDE · PLATFORM DOCS · CHECKED 4 NOV 2026

cited_text: “Up to 150 characters of the cited content”

PERPLEXITY · AGENT API PRESETS · OCT 2026

“web_search max_tokens and max_tokens_per_page rise from 2,000 to 4,000”

XAI · SEARCH TOOLS DOCS · CHECKED 5 NOV 2026

allowed_domains: “Only search within specific domains (max 5)”

XAI · CITATIONS DOCS · CHECKED 5 NOV 2026

“not every URL in this list will necessarily be directly referenced in the final answer”

Developer documentation, quoted as written. Not how consumer apps choose citations.

READS MORE

Engines read more of your page than they quote

PERPLEXITY · BEFORE

2,000

web_search max_tokens and max_tokens_per_page

PERPLEXITY · OCT 2026

4,000

Same settings, low, medium and high presets

Read: up to 4,000 tokens per page (Perplexity). Quoted: up to 150 characters (Claude). The quote is the short part.

Perplexity API changelog, Oct 2026 (Agent API presets). Claude Platform Docs, checked 4 Nov 2026. Different engines and units.

SEEN IS NOT CITED

A source it saw is not always a source it cites

> **“…not every URL in this list will necessarily be directly referenced in the final answer.”**

xAI docs, on the sources list returned with an answer. Retrieved and cited are two lists: measure both.

Ranking can get you retrieved. A short, self-contained claim gives it something to quote.

xAI docs, citations page, checked 5 Nov 2026. Last line: my read.

WHAT I WOULD DO

This week, 15 minutes

1

10 min · Rewrite the key claims

Pick your top page. One sentence per section, under 150 characters, with the number and the source inside.

2

2 min · Answer on the first screen

Put the answer where the page starts, not after the intro.

3

3 min · Start a sheet

Same prompts every month. Log retrieved vs cited, engine by engine.

Retrieved vs cited, engine by engine

PROMPT	ENGINE	RETRIEVED	CITED
Your prompt 1	ChatGPT	yes / no	**yes / no**
Your prompt 1	Perplexity	yes / no	**yes / no**
Your prompt 1	Claude	yes / no	**yes / no**
Your prompt 1	Grok	yes / no	**yes / no**
Your prompt 2	...	...	...

Retrieved: your URL is in the sources the engine shows. Cited: the answer quotes or links your page. One row per prompt and engine, same date each month.

A template, not data. Every engine sets different limits, so compare each engine with itself over time.

A NOTE ON THE SOURCES

What these docs do not show

These are developer docs for the APIs, not how the consumer apps choose citations. The Claude and xAI pages show no date, so I give the day I checked them.

Which page on your site would you rewrite first, and why?

Sources: Claude Platform Docs, web search tool (checked 4 Nov 2026). Perplexity API changelog (Oct 2026). xAI documentation, search tools and citations (checked 5 Nov 2026). Links in the first comment.