How well do AI chatbots answer news from today?

4 of 6

chatbots scored above 90% on multiple choice about events reported hours earlier

Now the fine print.

Preprint by Stanford University and Together AI, 21 May 2026.

THE STUDY

Who tested what, and when

WHO

Researchers at Stanford University and Together AI

STATUS

arXiv preprint, 21 May 2026. Not confirmed as peer reviewed

TESTED

Six chatbots with web search on: GPT-5, GPT-4o mini, Gemini 3 Pro, Gemini 3 Flash, Grok 4, Claude 4.5 Sonnet

QUESTIONS

2,100 multiple choice questions generated from same-day BBC News articles in six regional services

WHEN

14 days, 9 to 22 February 2026

Sources: arXiv 2605.22785; Stanford HAI summary, 3 June 2026.

Open answers score lower

MULTIPLE CHOICE

86.9%

Accuracy on a separate one-day check, 6 April 2026 (850 paired items)

OPEN ANSWERS

70–71%

Same questions, graded by three AI judges

The authors say multiple choice likely inflates accuracy compared with real use.

Across the 14 days in February, the four best chatbots passed 90% on multiple choice.

What happens without search

31–46

points more errors with web search turned off

US and Canada questions.

Where errors came from: retrieval failure and source divergence made up 71.5% of the 1,497 errors the authors classified. When the right source was retrieved, the model almost always answered correctly.

Source: arXiv 2605.22785; Stanford HAI, 3 June 2026.

Citing more is not answering better

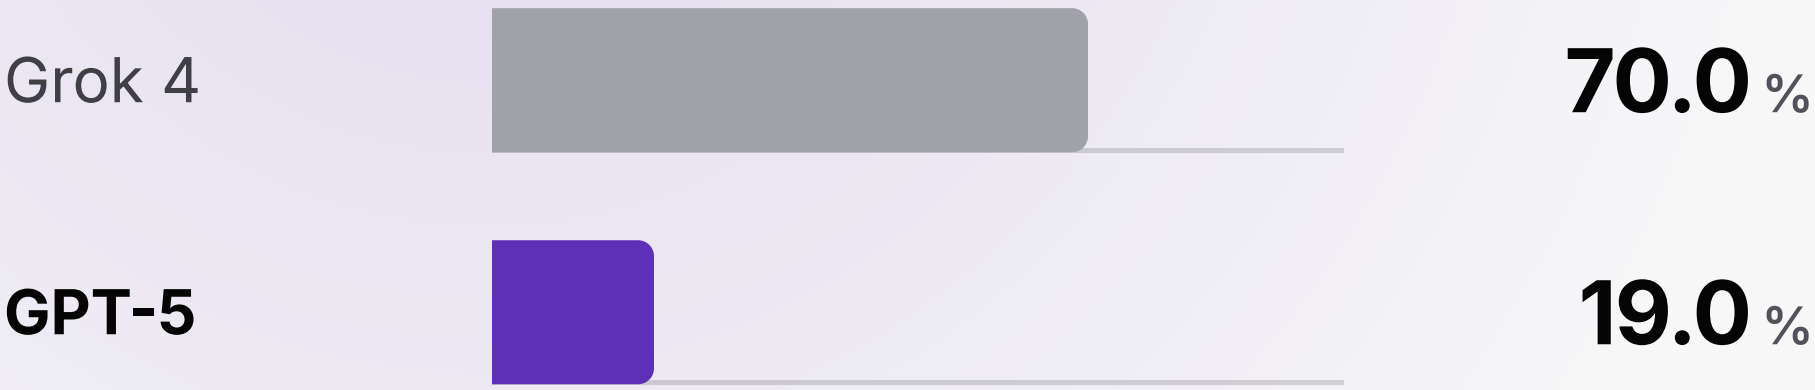
CHATBOT	ACCURACY	CITATION RATE
Gemini 3 Flash	**95.6%**	98.3%
Grok 4	**95.0%**	88.1%
Gemini 3 Pro	**93.7%**	84.7%
Claude 4.5 Sonnet	**90.4%**	96.2%
GPT-5	**85.0%**	97.8%
GPT-4o mini	**69.0%**	86.2%

The paper reports that citation frequency did not predict accuracy (ρ = 0.31, p = 0.54, six models).

Multiple choice accuracy. Citation rate: share of responses that cite at least one URL. Source: arXiv 2605.22785.

FALSE PREMISES

Hidden wrong assumptions were the weak spot



GPT-5 accepted false premises 64% of the time.

Four days, 100 items per model, US and Canada, English only. The authors call it suggestive, not definitive.

What the study does not show

STATUS
Preprint, not confirmed as peer reviewed

TIME
A 14-day snapshot from February 2026

SOURCE
BBC is unusually well indexed, so less prominent sources may do worse

FORMAT
Multiple choice likely inflates accuracy

QUESTIONS
All generated by one model, Gemini 3 Flash

LANGUAGES
No Spanish and no Mexico: English, Arabic, French, Hindi, Russian and Turkish

Source: arXiv 2605.22785; Stanford HAI, 3 June 2026.

WHAT I WOULD DO

This week, four steps

1. **Check my key pages are indexed**
   Search Console can tell you.

2. **Ask a chatbot with search on**
   My customer’s question. Save the answer, sources and date.

3. **Ask again with a wrong assumption**
   A feature I do not offer. Does it correct me?

4. **Repeat in a month**
   Answers change.

MY READ · OPINION

My read: retrieval first

In this study, most errors came from what the engine retrieved, not from how it phrased the answer.

Google’s documentation says a page must be indexed and eligible to show with a snippet to appear in its AI features. Whether answering in the first lines helps is something this study did not test.

When an AI answers about your business, do you know which sources it used?

Sources: Suzgun et al., “Evaluating Commercial AI Chatbots as News Intermediaries,” arXiv preprint (21 May 2026). Stanford HAI (3 June 2026). Google Search documentation on AI features.